

Io, gli Elefanti e la Teoria dei Giochi

University of Mannheim Department of Political Science

Formal Models in Comparative Politics & International Relations

Filippo Panfoli 2202955

✉ filippo.panfoli@students.uni-mannheim.de |  [PanfoliF](#) |

 0009-0000-7861-8036 |  <https://github.com/PanfoliF>

Replication materials / $\text{\LaTeX}$ sources:  [PanfoliF/Paper-Formal-Models](#)

Version: v1.1 | **Last updated:** December 31, 2025

Keywords: game theory; signaling; animal behavior; formal models; cooperation

JEL: C72; D82

Prepared for: Formal Models in Comparative Politics & International Relations,
taught by Prof. Dr. Thomas J. Bräuninger

License: CC BY 4.0

December 31, 2025

Abstract

Standard game-theoretic accounts often rationalize departures from fully rational play by appealing to exogenous "shocks" or moments of surprise. Yet many strategic encounters suggest that what matters is not merely an external disturbance but an endogenous shift that reshapes perception, signaling, and restraint. This paper develops an illustrative framework for endogenizing surprise and empathy in signaling settings, using two elephant-human encounters reported by Safina (2015) as motivating cases. In one episode, an elephant charges a tourist yet brakes before contact; in another an elephant injures a shepherd and subsequently guards him until help arrives. Canonical signaling models can capture incentives and credible threats in both interactions, but they struggle to accommodate the joint pattern "attack, then care." I treat the two episodes as potential outcomes of a common underlying interaction, and use a causal directed acyclic graph to make explicit the missing variables that standard models implicitly omit. I argue that surprise operates as an informational shock that can trigger aggression, while empathy functions as an affective regulator that modulates signal intensity and can generate post-conflict assistance. To formalize this intuition, I sketch an adaptation of Groseclose and Snyder's (1996) vote-buying game, in which competing emotions allocate limited intensity across competing "impulses" to determine behavior. The resulting perspective offers a tractable route to richer models of bounded rationality and yields testable implications about when shocks produce escalation versus restraint.

1 Introduction

In this paper I wish to describe an intuition regarding the way game theory has yet to formalize certain interactions between agents. Game theorists consider surprises and shocks as the rationale for an irrational behaviour. I believe that the following paragraphs will convince the audience of the value of this contribution for two reasons; first, I hope the description provided will be used and applied further. Second, the intuition about the role of emotions adds a new perspective to the literature about bounded rationality.

The intuition concerns animal behaviour that has drawn a lot from game theory. Even though it may appear far from what economists study, still, it deals with the problems that game theorists often face: the first is to represent unobserved mental processes, this means taking into consideration and inferring unobserved variables. The second problem deals with the creation of a formal idea about these processes. The last problem connects the formal idea to observed behaviour and derives an equilibrium, thus a state of the world with particular properties. I aim to show that seeking an explanation outside the mind is unnecessary, since it may well lie within it.

In order to ease the description and to make the model more directly applicable, I exploit a singular situation described in Safina (2015). This book is an amazing challenge for every person who has ever had to face a game-theoretic exercise. To support his thesis, the author provides the reader with detailed descriptions of animal interactions. All species selected belong to the realm of "clever" animals, animals that are capable of emotions and complex cognitive processes. I believe that each of these interactions is a useful game-theoretic exercise, since they force game theorist to apply its knowledge out of his area of expertise.

This book also provides a straightforward setting for these exercises, since animals provide a simplified set of agents for game theory. They are complex enough to be represented as rational players, but they are not as complex as humans, so that most of the assumptions and simplifications have more grasp into reality. In this respect, elephants appear to be the perfect application. There is plenty of evidence about their behaviours, and there is plenty of evidence that they act in accordance with a *mind*. According to Safina (2015), they are almost humans, and they often interact with humans.

I will formalise my intuition through a Causal Directed Acyclic Graph (DAG). Since I wish to narrow down the gap between theory and practice, I believe this tool is well suited to the research design. Additionally, it provides a more transparent version of the analysis, clarifying the variables of interest and the limits of the assumptions.

The final aim of this article is to clarify a contradiction between two similar situations that game theory formalises in different ways.

1.1 Beyond Words

Carl Safina is a biologist and has written a widely read essay (Safina, 2015) about animals' minds. He argues that until now, biologists have been reluctant to deepen the study of animals' minds and "souls" and that is high time they did so. The curious thing about this book is that, in doing so, he draws readers' attention a problem quite close to game theory. He tries to rationalise the behaviour of some animals, arguing that their behaviour may be explained if we accept they are "intelligent", meaning that a mind guides their behaviour and that they are sensitive to emotions. This book is an attempt to place animals' minds at the same level of human minds. **I hope the animal-world examples will:**

- change the way we theorise a broader set of agents and interactions.
- be used to test and reinforce the empirical evidence.
- motivate a refined equilibrium concept.

1.2 The Problem

This section aims to clarify the setting captured by the game-theoretic models, as well as making the reader aware of the basic intuition. In the examples Safina describes in his book (Safina, 2015) animals seem to be driven by rationality and emotions, much like humans. *These are brief summaries of the two episodes.* Both interactions between humans and elephants happened in a wildlife park in Africa. In both, the surrounding environment appears to be similar.

Case 1:

A staff member of a wildlife park in Africa reports an "attack" by an elephant on a tourist. The scene clearly highlights the signals that the elephant sends to the tourist. It also clearly identifies the reasons why it is logical to believe that the elephant intentionally avoided hurting the tourist. In particular, the staff member reports signs on the ground that undoubtedly mark the intention to halt the charge against the tourist moments before impact. **We can formalise this game through a signalling game in which the elephant sends a threatening signal without intending to harm him.**

Another situation, recorded in the same environment, shows how similar conditions can produce an opposite pattern of behaviour.

Case 2:

In the same wildlife park in Africa, a rescue group was involved in a mission to reach a shepherd who was badly hurt by an elephant. When the staff approached him, they reported that the shepherd stopped them from harming the animal. He argued the elephant had kept him safe during the night. From the behaviour of the animal it seemed clear what it intended. After the attack, it soon realised it had misinterpreted the signals from the shepherd, and after hurting him, the animal

took care of him until the rescue group reached the place of the accident. Here, the same signalling logic as in the first example breaks down: the equilibrium is reached, yet the elephant attacks. However, the event suggests a richer model of rationality, where the animal is able to recognise an "error".

Both scenarios can be framed within game theory, yet each reveals elements of cognition and emotion that standard models overlook.

In particular, we shall consider three puzzles that emerge from these examples:

1. Why does the animal attacks and then provides assistance? Isn't it contradictory?
2. Why does the animal initiates attack only in one of the two cases?
3. What is rational? To attack and regret the decision, or not attack at all? With respect to what criteria do we define which is the best strategy?
4. Which is the best strategy in evolutionary terms?

1.3 Parallel with Game Theory

I believe game-theoretic analysis accurately captures the key strategic features of both cases. For example, the fact that the animal starts a conflict with humans and tends to be pacific with most of other creatures is due to the fact that humans are one of the few threats they face in nature. Game theory correctly models agents as driven by incentives, yet it does not take into account interactions in which the agent acts violently because of *its malicious nature*. Similarly, game theory models credible threats as equilibria depending on the information available to players: not the tourist nor the sheperd intended to intentionally harm an animal that could plausibly kill them, for example.

Additionally, there are plenty of models that suite each of the two situations. For instance, as we will see, some models explain the behaviuor of the first case thorough the concept of mimicry: the elephant knows it will not cause harm to the tourist. However, the tourist cannot distinguish between an animal with good or bad intentions, and the elephant exploits this weakness for defense porposes.

Thus, both animals and humans may be depicted as rational agents, since their behaviour appears to be consistent with the incentives they face and with a broader goal, if not with a specific aim in mind. Both appear to be able to communicate with each other, indicating that both species are adapting to signals from other agents and that they can act as informed agents. And both seem to be aware of each other's strengths, at least partially (Schelling, 1960; J. M. Smith, 1973). Again, in both cases, the animal seems "pacific" and willing to avoid or prevent conflict, at least up to a point.

1.4 Representative Agents

In both cases analogies between humans' and animals' minds do not appear to be straightforward. However, most of their differences need not be represented

in game theory formalisations. Additionally, game theory is often applied at different levels of analysis (e.g. nations' behaviour), which means that only the crucial features of agents need to be represented.

Thus, under the descriptive limits of these games and under their assumptions, I believe no confusion arises when I use "human", "elephant", and "agent" interchangeably. *I will differentiate between these agents only in those cases in which it is necessary, and with due warning to the reader.*

2 Causal Inference

The discussion so far motivates a comparison with frameworks that model how unobserved factors influence observed outcomes. Causal inference (Pearl, 2009) suggests that in order to answer this question, we may also think of an unobserved confounder that is mistakenly excluded from the framework. I argue that an omitted variable is present in this setting.

2.1 Potential Outcomes Framework

This section formalises the intuition using a DAG. To make the subsequent discussion accessible, it first outlines how the potential-outcomes framework applies to this setting.

A potential outcome is a counterfactual. Thus, it is a scenario that we do not observe; we think of it as the closest and most realistic scenario that might have happened had the actual one not occurred.

In the case of the two interactions I presented, each of them may be considered the counterfactual scenario of the other. Since counterfactual outcomes do not exist, we use real and observed cases to mimic them. The underlying belief is that no researcher is as informative as reality itself in producing examples of what might have happened. If we rely on observation, the role of the researcher is limited to selecting the most appropriate counterfactual from those offered by the natural world.

Hypothesis: in the absence of emotional effects, the case involving the elephant and the shepherd would have been identical to the one offered by the other elephant and the tourist.

2.2 Application of the Potential Outcomes Framework




Figura 1: *Causal DAG.*

The graph shows five circles that represent our variables. We temporarily set aside the emotion node to describe the baseline case.

In the baseline setting agents are represented by perceptions, rationality and behaviour. The former two shape the latter, which is the outcome variable of interest. A change in perceptions corresponds to a change in behaviour. Rationality is modelled as the decision rule or the algorithm agents have in their minds. It is the function that represents utility. Two players with the same utility function act in the same way when facing the same game. This means that a change in rationality corresponds to a change in behaviour.

Our outcome variable is the only one we can observe and measure. In the two cases described by Safina (2015), we *observe* the behaviour of elephants and humans. However, we can only *infer* rationality and perceptions.

When we represent two agents using three variables, we make use of three circles for each player. Nonetheless, the figure presents only one circle for perceptions and only one circle for rationality.

This choice reflects a simplification. In order to ease the analysis, we assume that both agents perceive reality in the same way. Thus, both receive the same stimuli from the environment. We also assume that they process these stimuli through the same algorithm. They weigh the same benefits and costs in the same way. These assumptions simplify the baseline setting because restrict the set of potential outcomes: we consider only two possible outcomes: rationality in state

A or B and perceptions in state A or B .

How do we distinguish between the two agents in the Causal DAG? They differ in the realised behavioural outcome. While human behaviour is constant, elephant is different in the two cases (The Problem), this is why we have one circle for each player.

We have 4 circles, three of which are constant in the two cases. The fifth one represents elephant's emotions. This variable provides the key source of variation across cases. Thus, it can explain the change in elephant behaviour.

2.3 The Rationale for using Causal Inference From the General Case to a Single Case—and Back Again

von Neumann and Morgenstern (1990, chap. 2, p. 14) write: "An almost exact theory of a gas, containing about 10^{25} freely moving particles, is incomparably easier than that of the solar system, made up of 9 major bodies; and still more than that of a multiple star of three or four objects of about the same size. This is, of course, due to the excellent possibility of applying the laws of statistics and probabilities in the first case. [...] The problem must be formulated, solved and understood for small numbers of participants before anything can be proved about the changes of its character in any limiting case of large numbers".

To pursue this suggestion, we go move the abstract general case of rational-agent model to analyse specific interactions in which the model fails to adequately map reality.

We then turn to causal inference in order to reconstruct a more general account: namely, to move from the idiosyncrasies of a single interaction to a *new* general framework. This requires abstracting away from irrelevant features and identifying a shared causal mechanism. The result is a revised pattern: a new map that can guide inference, organize the evidence, capture the data-generating process.

3 Models

We restate the three puzzles that emerged from the animal examples in subsection 1.2 (The Problem):

1. Why does the animal attacks and then provides assistance? Isn't it contradictory?
2. Why does the animal attacks only in one of the two cases?
3. What is rational? To attack and regret that decision or not attack at all? With respect to what criteria do we define which is the best strategy?
4. Which is the best strategy in evolutionary terms?

3.1 Model 1 - Surprise as Shock

Reiter (2003) argues: "Classic deterrence theory predicts that as an imbalance of power between two states grows, war becomes more probable. The bargaining model, *on the other hand*, argues that it is not an imbalance of power that causes war, but rather disagreement over the balance of power (and over the likely outcome of war)." (*Emphasis added*).

In the light of formal models of conflict theory, the first question can be rephrased as following: why does the animal acts through deterrence in one case and through bargaining in the other?

At first glance, we may theorise that the animal regrets the attack

The decisive variable in this situation may be the element of surprise (Fearon, 1995). The two versions of the elephant-human encounter, one involving the shepherd and one involving the tourist, are similar in structure. In the tourist scenario, there is no surprise: the elephant is aware that humans might approach and disturb it, and it knows where to find them.

The situation involving the shepherd, however, contains the element of surprise, which may justify a different behavioural response.

3.2 Model 2 - Vote Buying

3.2.1 Back to Groseclose and Snyder (1996)

In this section I propose a reworking of a well-known model by Groseclose and Snyder (1996). I intend to build on the equilibrium the authors derive, to support the intuition developed in the DAG with another narrative. In doing so, I aim to improve on the previous model (See subsection 3.1: Model 1 - Surprise as Shock).

In their article, Groseclose and Snyder (1996) explain why, in the US Congress, the number of votes often exceeds the minimum required to make a bill pass. They argued that this is due to the interaction between lobbying groups competing for support on different bills. These groups *bribe* Congressmen to convince them.

The basic intuition from the paper comes from this example:

"The legislature contains seven members, all indifferent between the two proposals, that is, $v(i) = 0$ for all i . Since B will move second and attack the weakest part of A 's coalition, A will offer the same bribe, a , to all legislators he bribes. Let $m + 4 \geq 0$ be the size of the coalition A bribes. If $m = 0$, then B can defeat x by paying $a + \varepsilon$ to one member of A 's coalition. Thus, A must set $a \geq W_B$ to win, and A 's total bribes are then $4W_B$. If $m = 1$, then B must pay $a + \varepsilon$ to two members of A 's coalition to defeat x . Thus, A must set $a \geq W_B/2$ to win, and his total bribes are then $5W_B/2$. This is less than $4W_B$, so A is better off buying a supermajority of size 5 than a minimal winning coalition. If $m = 2$, then B must pay $a + \varepsilon$ to three members of A 's coalition, so A must set $a \geq W_B/3$ to win, and

his total bribes are then $2W_B$. In fact, A minimizes his total payments by bribing all seven legislators. Here, $m = 3$, and A 's total bribes are $7W_B/4$. Clearly, there is an equilibrium where x wins if and only if $W_A \geq 7W_B/4$. If $W_A < 7W_B/4$, then there are only equilibria in which no bribes are paid and s wins." (Groseclose & Snyder, 1996)

3.2.2 Representation through Vote Buying

Consider a council composed of several members. These members are not-strategic players whose votes can be influenced. There are two active players, conceptualized as two groups, each seeking to "buy" the votes of the council members in order to determine the policy agenda. Each group possesses a budget, and bills are voted on by the council according to a pre-defined rule, assumed here to be the simple majority rule. This setup corresponds to the classical *vote-buying game* as formalized by (Groseclose & Snyder, 1996).

The objective of this analysis is to adapt this framework to the so-called *elephant-shaped* case. This adaptation allows the original model's equilibrium structure to be embedded in a cognitive-emotional setting.

- Let K denote the total number of *thoughts*.
- The two players represent *opposing emotions* that govern the internal dynamics of the elephant's cognition. "Setting the agenda" refers to determining the behaviour.
- Each player's *budget* corresponds to the *intensity* of its emotion. A higher emotional intensity enables the emotion to "purchase" a larger number of thoughts.

Notation:

Players: $N = \{X, Y\}$ or $N = \{Emotion1, Emotion2\}$.

Histories: $H = \{\emptyset, \{x\}, \{x, y\}\}$. Where x and y indicate payments vectors in the baseline setting. Emotions use units of payment (emotional intensity) to corrupt a thought (or a judgement).

Player function: $p(\emptyset) = X$; $p(x) = Y$

Actions: $A(\emptyset) = \{x\}$; $A(x) = \{y\} \forall x$.

Each player *receives* V_i if elephant acts accordingly.

Each player loses $\sum_{j=1}^k x_j$ for each unit of intensity spent.

This adapted framework enables a formal representation of conflict as a strategic interaction between emotions competing for cognitive control. Each emotion allocates its limited intensity across thoughts, attempting to secure a majority that is sufficient to determine the elephant's behavioural choice. The resulting

equilibrium can be interpreted as a stable emotional configuration in which neither emotion can improve its outcome given the other's allocation strategy. Let μ denote the minimal number of thoughts required to *set the agenda*, that is, to achieve behavioural dominance of one emotion over the other. Then:

$$\mu = \frac{k + 1}{2}$$

4 Literature Review and Limitations

4.1 A Review of Models of Conflicts

Further exploration of the conflict-games literature either support or weaken the analysis presented so far. In any case, already-established models provide a useful point of comparison. Some of the papers that stand out include (Schelling, 1960, 1966). On the other hand, it may be of interest to focus on McKelvey and Palfrey (1995) to further develop the model.

4.1.1 A Review of Models of Conflicts with a Focus on Bargaining

A specific branch of the literature focuses on bargaining models of conflicts. Key contributions to this field include Fearon (1995, 1996, 1998) and Schelling (1960, 1966). Reiter (2003) offers an in-depth review of the field. He provides an historical perspective and discusses more recent works. Additionally, he summarises the most important ideas.

Reiter (2003) says: "James Fearon prominently contributed to this scholarship with his 1995 paper "Rationalist Explanations for War." He usefully highlighted the point that if states agreed on the outcome of a possible war, they could probably avoid war. Central to Fearon's paper is the importance of focusing on states' disagreement over their capabilities and/or resolve. [...] Fearon uses a bargaining model to argue that if two states in dispute know the outcome of a possible war, they should in general prefer to reach a deal that would reflect the hypothetical postwar political settlement, rather than fight, reach that same settlement, and also suffer the costs of war. This view assumes that fighting itself is costly—that the belligerent always suffers some negative utility, no matter how the issue at stake is settled. In fact, this assumption becomes necessary when one asserts that two states would always prefer to reach a bargain without fighting rather than fight and then reach the same bargain. Within his bargaining model, Fearon develops three conditions under which war is possible."

Modelli di cherim: In order to broaden the overview, some other interesting contributions come from Filson and Werner (2004), Kirshner (2000), A. Smith and Stam (2004), Wagner (2000), and Wittman (1979).

4.2 Limits of Theoretical Models

This discussion is necessarily selective and cannot provide a comprehensive assessment of the formal-modelling literature. In the previous pages I have presented

my personal re-elaboration along with a brief summary of the relevant literature I studied in this area. This is certainly not enough to assess the state of the models of this article. However, the exposition I provide the reader with, should not be considered sufficient in the sense that there is no such thing as *a single correct or definitive modelling framework*.

In this respect, it is possible to further expand the discussion on limitations. The literature offers some very good papers, among others, I discuss: Ashworth et al. (2021), Bräuninger and Swalve (2020), and Pareto (1916). The present different ways to combine empirical and theoretical research together with some insightful interpretations of the meaning and scope of models.

4.3 Connection between Game Theory and Biology

There are many important applications of game theory to biology. I believe the most interesting ones are: Axelrod and Hamilton (1981) and J. M. Smith (1973). Future work may shed light on more recent, as well as interesting applications. A vast literature covers the topic of evolutionary games too: Hofbauer and Sigmund (2003), Kandori et al. (1993), Nowak (2006), and Taylor and Jonker (1978). In particular, Hofbauer and Sigmund (2003, chap 12) offers a clear treatise of powerful analytical tools together with applications of these game-theoretic tools in biology.

Unfortunately, a more detailed discussion of these works must be postponed.

Conclusion

The puzzle motivating this paper was deliberately simple: two elephant-human encounters occur in broadly similar environments, yet they generate sharply different behavioural trajectories. In one case, the elephant's threat is forceful but calibrated; in the other, violence is followed by a striking form of care.

Standard signalling logic is largely successful in explaining why an elephant might issue a credible threat and why a human might respond cautiously. What it does not readily explain, is the combination of escalation and subsequent protection—an outcome that looks, from the standpoint of canonical rational-agent models. The central claim advanced here is that the apparent inconsistency largely reflects omitted elements of the model: the relevant "shock" can be internal, and the relevant constraint can be affective.

The causal-inference lens is useful. Treating the two episodes as potential outcomes of a shared underlying interaction makes the research design explicit: the cases serve as mutual counterfactuals, and the goal is to isolate which latent factor, omitted in standard game-theoretic models, generates the divergence. Within this framing, surprise functions as an informational discontinuity that plausibly shifts the elephant's perception of threat and thus the equilibrium-relevant mapping from signals to actions. Crucially, however, surprise alone cannot account for the post-attack behaviour. That second step motivates the introduction of empa-

thy as a regulator: a force that can attenuate aggression, reshape the "intensity" of signals, and generate assistance once the situation is reinterpreted.

This conceptual move has a modelling payoff. Instead of treating emotions as noise perturbs otherwise coherent strategic choice, the paper sketches a way to represent emotions as structured, competing influences over cognition. The adaptation of the vote-buying framework provides a particularly transparent metaphor: thoughts play the role of passive "voters", while emotions act as active players allocating scarce intensity to secure agenda-setting power—here, the determination of behavior.

In equilibrium, the elephant's observable trajectory corresponds to a stable emotional configuration: neither emotion can improve its outcome given the other's allocation. Read in this way, "attack, then care" is not a contradiction so much as the **observable trace of a reallocation of cognitive control after misinterpretation is recognized**.

The broader implication is that many phenomena labeled as "shocks" in political science and international relations may be mischaracterised when treated as purely exogenous disturbances. A framework that makes these mechanisms explicit can therefore complement familiar strategic models, not by abandoning rationality, but by widening the set of state variables that rational choice is allowed to condition on.

At the same time, the limits of the present contribution should be stated plainly. A fuller treatment would require specifying payoffs, information structures, and equilibrium selection more rigorously. Future work could therefore proceed by tightening the formal model, including the relationship between surprise and empathy as distinct but interacting state variables.

The main takeaway message is modest: when game theory explains puzzling behaviour by invoking "surprise", some times it may be pointing towards a missing internal mechanism. Making that mechanism explicit—rather than treating it as residual noise—can turn an ad hoc rationalization into a tractable extension of formal modeling, one that is better aligned with the texture of real strategic behaviour.

Acknowledgements

Artificial intelligence-based tools were employed solely to improve linguistic clarity and grammar. No AI system contributed to the development of the research questions, theoretical framework or conclusions presented in this paper.

Bibliography

Books

- Ashworth, S., Berry, C. R., & Bueno de Mesquita, E. (2021). *Theory and credibility: Integrating theoretical and empirical social science*. Princeton University press.
- Hofbauer, J., & Sigmund, K. (2003). *Evolutionary games and population dynamics* (Transferred to digital printing). Cambridge University Press.
- Pearl, J. (2009). *Causality: Models, Reasoning, and Inference* (2nd ed.) [Chapters 1 and 2 define causal diagrams (DAGs) and the structural model framework.]. Cambridge University Press.
- Safina, C. (2015). *Beyond Words: What Animals Think and Feel* [Explores animal cognition and emotions, linking behavioral evidence with questions of consciousness.]. Henry Holt and Company.
- Schelling, T. C. (1960). *The Strategy of Conflict* [Classic work on strategic behavior and conflict theory.]. Harvard University Press.
- Schelling, T. C. (1966). *Arms and influence* [Classic work on deterrence, coercion, and the strategic use of uncertainty in conflict.]. Yale University Press.
- von Neumann, J., & Morgenstern, O. (1990). *Theory of games and economic behavior* (6., paperback printing). Princeton University Press.

Articles

- Axelrod, R., & Hamilton, W. D. (1981). The Evolution of Cooperation.
- Fearon, J. D. (1995). Rationalist explanations for war. *International Organization*, 49(3), 379–414. <https://doi.org/10.1017/S0020818300033324>
- Fearon, J. D. (1998). Bargaining, enforcement, and international cooperation. *International Organization*, 52(2), 269–305. <https://doi.org/10.1162/002081898753162820>
- Filson, D., & Werner, S. (2004). Bargaining and fighting: The impact of regime type on war onset, duration, and outcomes. *American Journal of Political Science*, 48(2), 296–313. <https://doi.org/10.1111/j.0092-5853.2004.00071.x>
- Groseclose, T., & Snyder, J. M. (1996). Buying Supermajorities (2014/08/01). *American Political Science Review*, 90(2), 303–315. <https://doi.org/10.2307/2082886>
- Kandori, M., Mailath, G. J., & Rob, R. (1993). Learning, Mutation, and Long Run Equilibria in Games. *Econometrica*, 61(1), 29. <https://doi.org/10.2307/2951777>
- Kirshner, J. (2000). Rationalist explanations for war? *Security Studies*, 10(1), 143–150. <https://doi.org/10.1080/09636410008429423>
- McKelvey, R. D., & Palfrey, T. R. (1995). Quantal Response Equilibria for Normal Form Games. *Games and Economic Behavior*, 10(1), 6–38. <https://doi.org/10.1006/game.1995.1023>
- Nowak, M. A. (2006). Five rules for the evolution of cooperation. *Science*, 314(5805), 1560–1563. <https://doi.org/10.1126/science.1133755>

-
- Reiter, D. (2003). Exploring the Bargaining Model of War. *Perspective on Politics*, 1(1), 27–43. <https://doi.org/10.1017/S1537592703000033>
- Smith, A., & Stam, A. C. (2004). Bargaining and the nature of war. *Journal of Conflict Resolution*, 48(6), 783–813. <https://doi.org/10.1177/0022002704268026>
- Smith, J. M. (1973). The logic of Animal Conflict. 246.
- Taylor, P. D., & Jonker, L. B. (1978). Evolutionary stable strategies and game dynamics. *Mathematical Biosciences*, 40(1–2), 145–156. [https://doi.org/10.1016/0025-5564\(78\)90077-9](https://doi.org/10.1016/0025-5564(78)90077-9)
- Wagner, R. H. (2000). Bargaining and War. *American Journal of Political Science*, 44(3), 469–484. <https://doi.org/10.2307/2669259>
- Wittman, D. (1979). How a war ends: A rational model approach. *Journal of Conflict Resolution*, 23(4), 743–763. <https://doi.org/10.1177/002200277902300408>